

P7: Design an A/B Test

The full version of the project: https://olgabelitskaya.github.io/P7_Design_an_A_B_Test_Overview.html

Experiment Design

Metric Choice

Invariant metrics (expected to be unchanged in the control and experimental groups):

- 1) number of cookies (cannot be affected by the experiment: users made a decision to visit the page before they were asked the question);
- 2) number of clicks (cannot be affected by the experiment: users clicked the button before they were asked the question);
- 3) click-through probability (cannot be affected by the experiment: it equals to the number of clicks divided by the number of cookies).

Evaluation metrics (expected to be different in the control and experimental groups):

- 1) gross conversion (can be affected by the experiment / can decrease: users could make a decision to enroll in the free trial in the experimental group less than in the control group because they did not plan to learn 5+ hours per week);
- 2) retention (can be affected by the experiment / can increase: enrolled users could be disappointed in the learning process less and make more payments in the experimental group than in the control group because they paid attention to studying 5+ hours per week);
- 3) net conversion (can be affected by the experiment / can decrease: users could enroll in the free trial less in the experimental group than in the control group, thus could decrease the number of people who paid).

The goals of the experiment in the practical meaning:

- the number of payments should not be decreased;
- the number of students who were disappointed and had not paid because they could not study enough time should be reduced.

The goals of the experiment in terms of our metrics:

- the gross conversion should significantly decrease;
- the retention should significantly increase;
- the net conversion should not decrease.

An important remark: the number of user-ids is neither a good invariant metric nor a good evaluation metric.

From one side, the new pop-up message is likely to decrease the total number of user-ids who enrolled in the free trial, so it is not invariant; from the other side it is not normalized, the number of visitors may be different between the experiment and control groups, so it is not good for evaluation.

Measuring Standard Deviation

Number of cookies = 5000

Number of clicks on "Start free trial" = $5000 \times 0.08 = 400$

Number of enrollments = $5000 \times 0.08 \times 0.20625 = 82.5$

$$SD \text{ Gross conversion} = \sqrt{\frac{p*(1-p)}{n}} = \sqrt{\frac{0.20625*(1-0.20625)}{400}} = 0.0202$$

$$SD \text{ Retention} = \sqrt{\frac{p*(1-p)}{n}} = \sqrt{\frac{0.53*(1-0.53)}{82.5}} = 0.0549$$

$$SD \text{ Net conversion} = \sqrt{\frac{p*(1-p)}{n}} = \sqrt{\frac{0.1093125*(1-0.1093125)}{400}} = 0.0156$$

I would like to expect the analytical variance is close to the empirical variance for the gross conversion and for the net conversion: the denominator for these two indicators is the number of clicks, which is also the unit of diversion.

And it would be useful to collect an empirical estimate of the variability for the retention: the unit of diversion was not used in this case, the empirical variance of the retention is more likely to be higher than the analytical variance.

Sizing

Number of Samples vs. Power

I have used the online calculator (References, N5) for calculating the sample sizes and have chosen the largest.

Gross conversion: $2 \times 25835 \times 40000 \div 3200 = 645875$ pageviews

Retention: $2 \times 39115 \times 40000 \div 660 = 4741212$ pageviews

Net conversion: $2 \times 27413 \times 40000 \div 3200 = 685325$ pageviews

I did not use the Bonferroni correction. Number of pageviews: 685325.

Duration vs. Exposure

Number of pageviews: 4741212 (for the experiment with the gross conversion, the retention and the net conversion)

Days (100% of the traffic) = $4741212 \div 40000 = 118.5303$

Number of pageviews: 685325

(for the experiment with the gross conversion and the net conversion)

Days (100% of the traffic) = $685325 \div 40000 = 17.133125$

Days (60% of the traffic) = $685325 \div 40000 \div 0.6 = 28.5552083333$

Number of pageviews: 685325 Fraction of traffic exposed: 0.6 Length of experiment: 29

Even the use of 100 percent of the traffic does not allow testing in full for all three metrics (the gross conversion, the retention, and the net conversion). It turns out that we need to experiment 118 days. Of course, it is too long for tasks staged in the project.

Percent reduction in traffic will increase this interval. Hence, it is necessary to reduce the number of test metrics. I should reconsider an earlier decision and setup the number of pageviews equal to 685325. For the experiment with the gross conversion and the net conversion, we can use the period 17 days with the 100% level. This interval is much better, but it gives us the result too quickly. The behavior of people in the field of education is quite difficult to analyze for 2-3 weeks. To slightly increase the time interval, we will set the percentage of used traffic at 60 (fraction = 0.6) and it gives us the number: 29 days.

However, I want to analyze the experimental data on my own initiative a little bit wider and to determine what happens to the metrics "Retention" as the most interesting in terms of the psychological effect and the most unpredictable indicator. I hope it is expedient to keep 100 percent of the traffic using all three metrics for analysis and it is a possible scenario in the future research because the audience of this site is very different in education, age, nationality, and other characteristics. Any reduction in the percentage of participants can significantly distort the results.

It is also possible to take into consideration that the commercial risk is low: the site offers students to better plan and to evaluate their training time. For investments, the decisive factor is the first free period as a way to assess their own abilities, and it does not change during the experiment. And there is no risk in terms of privacy violation at all. Users do not enter any additional information about themselves in the process of the experiment.

Experiment Analysis

Sanity Checks

Control group: Clicks = 28378 Pageviews = 345543

Experimental group: Clicks = 28325 Pageviews = 344660

Number of cookies:

Standard error SE = 0.000601840740294

Margin of error ME = 0.00117960785098

Confidential interval CI = (0.49882039214902313, 0.5011796078509769)

$\hat{p} = 0.500639666881 \in (0.49882039214902313, 0.5011796078509769)$ ✓

Number of clicks on "Start free trial":

Standard error SE = 0.0020997470797

Margin of error ME = 0.00411550427621

Confidential interval CI = (0.49588449572378945, 0.5041155042762105)

$\hat{p} = 0.500467347407 \in (0.49588449572378945, 0.5041155042762105)$ ✓

Click-through-probability on "Start free trial":

Pooled probability $p_{\text{pool}} = 0.0821540908979$

Standard error SE = 0.000661060815639

Margin of error ME = 0.00129567919865

Difference $\hat{d} = 5.66270915869e-05$

Confidential interval CI = (-0.0012956791986518956, 0.0012956791986518956)

$\hat{d} \in (-0.0012956791986518956, 0.0012956791986518956)$ ✓

Number of cookies

Lower bound = 0.4988; Upper bound = 0.5012; Observed = 0.5006; Passes = Yes

Number of clicks on "Start free trial"

Lower bound = 0.4959; Upper bound = 0.5041; Observed = 0.5005; Passes = Yes

Click-through-probability on "Start free trial" (Difference between the control and experimental groups)

Lower bound = -0.0013; Upper bound = 0.0013; Observed = 0.0001; Passes = Yes

All invariant metrics have stood the test successfully. This is a very predictable result. Selecting these metrics was based on the lack of experiment influence.

Result Analysis

Effect Size Tests

For each of our evaluation metrics, I gave a 95% confidence interval around the difference between the experiment and control groups and indicated whether each metric was statistically and practically significant.

Gross conversion

Pooled probability $p_{\text{pool}} = 0.208607067404$

Standard error SE = 0.00437167538523

Margin of error ME = 0.00856848375504

Difference $d = -0.0205548745804$

Confidential interval CI = (-0.0291233583354044, -0.01198639082531873)

$0 \notin (-0.0291233583354044, -0.01198639082531873)$

Retention

Pooled probability $p_{\text{pool}} = 0.551886792453$

Standard error SE = 0.0117297800914

Margin of error ME = 0.0229903689791

Difference $d = 0.0310948047071$

Confidential interval CI = (0.008104435728019967, 0.05408517368626556)

$0 \in (0.008104435728019967, 0.05408517368626556)$; $d_{\text{min}} = 0.01 \in (0.008104435728019967, 0.05408517368626556)$

Net conversion

Pooled probability $p_{\text{pool}} = 0.115127485312$

Standard error SE = 0.00343413351293

Margin of error ME = 0.00673090168535

Difference $d = -0.00487372267454$

Confidential interval CI = (-0.011604624359891718, 0.001857179010803383)

$0 \in (-0.011604624359891718, 0.001857179010803383)$; $d_{\text{min}} = -0.0075 \in (-0.011604624359891718, 0.001857179010803383)$

I did not use the Bonferroni correction.

Gross conversion (Difference between the control and experimental groups)

Lower bound = -0.0291; Upper bound = -0.0120; Statistical significance = ✓; Practical significance = ✓

Retention (Difference between the control and experimental groups)

Lower bound = 0.0081; Upper bound = 0.0541; Statistical significance = ✓; Practical significance = ✗

Net conversion (Difference between the control and experimental groups)

Lower bound = -0.0116; Upper bound = 0.0019; Statistical significance = ✗; Practical significance = ✗

Sign Tests

Gross conversion: success = 4 total = 23

Retention: success = 13 total = 23

Net conversion: success = 10 total = 23

I have used the online calculator (References, N6) for the sign tests.

I did not use the Bonferroni correction.

Gross conversion: p-value = 0.0026; Statistical significance = ✓

Retention: p-value = 0.6776; Statistical significance = ✗

Net conversion: p-value = 0.6776; Statistical significance = ✗

Gross conversion difference

Median: -0.0247583431053

Mean: -0.0207845820293

Retention difference

Median: 0.0232108317215

Mean: 0.0333425074664

Net conversion difference

Median: -0.00902785253995

Mean: -0.00489685698981

Summary

We have measured two (or three) metrics in one experiment. Applying the Bonferroni correction means that the α -level for each hypothesis will be 2.5 % or even 1.6% instead of 5% and confidential intervals will be significantly wider. It is too conservative, and I did not use this correction.

Our metrics have a strong relationship between each other. If we know the outcome of one test of a difference between the control and experimental groups on one metrics, it would be easy to predict and to find the outcome of the other tests on the other metrics. Their results will be more likely to occur not by chance and will have a strong relationship as well.

Eventually, the effective size and sign tests show that the site change would statistically significantly reduce the gross conversion, but would not affect the net conversion in a statistically significant way. The effect size test states this in the practical meaning also.

Statistical significance of the difference between the control and experimental groups was checked by using the values of the mean (effect size tests) and the median (sign tests). As we can see, these values vary considerably for the retention and this is the cause that the test results also differ.

Recommendation

The recommendation is not to launch the experiment change because the negative results outweighed the positive ones.

Positive results of the experiment.

- The difference for the gross conversion is practically significant and negative. This is a good sign: the Udacity team can lower costs by a number of trial signups.
- The difference for the net conversion is not statistically significant. It means the absence of serious financial losses.

Negative results of the experiment:

- The interval for the difference in the case of the net conversion includes negative numbers. Therefore, the team has a risk to decrease incomes.
- We have not gathered enough data to draw conclusions about the retention and because of this we can not evaluate correctly the difference between the control and experimental groups for a number of students who were disappointed in studying during the free period. Consequently, we do not know enrolled users would be disappointed in the learning process less and make more payments or would not.

If it is possible to continue the research in the similar format, it is necessary to improve the design of visualization for successful studying, increase the duration of the experiment for better understanding the behavior of students and expand the audience to 100 percent for avoiding bias. Perhaps this will detect certain trends in the metric "Retention" and decrease the financial risk.

Follow-Up Experiment

I would suggest these possible changes to the proposed experiment:

- extend the duration of the experiment up to 2 months with constant monitoring of incomes for avoiding financial risks;
- measure all three evaluation metrics for the 100%-audience;
- replace the visualization message by the video with an explanation of successful learning strategies based on statistics of the particular site or by the input test for the course level recommendations for this user.

I think the effect will be more detectable.